Wonje Jeung

Email Scholar Website LinkedIn

Education

University of Michigan *Sept 2026- May 2030*
PhD in Computer Science (Expected)

- **James F. Lu World Fellows Program**: One of 4-6 recipients across all departments of the College of Engineering, with $5,000 merit stipend/year and $10,000 candidacy development stipend.
- Advisor: Lu Wang

Yonsei University *Sept 2024 – Aug 2026*
MS in Artificial Intelligence (Expected)

- GPA: 4.24/4.3
- **Coursework:** Information Theory, Statistical Pattern Recognition, Multimodal Deep Learning, Text Understanding, Large Scale Learning and Inference, Medical Imaging, Responsible AI
- Advisor: Albert No

BS in Computer Science *Mar 2020 – Aug 2024*

- GPA: 4.09/4.3
- **Coursework:** Machine Learning, Computer Vision, Deep Learning Theory and Practice, Signal Processing, Probability and Statistics, Computer Systems, Computer Architecture, Operating Systems, Big Data
- Advisor: Jonghyun Choi

Publications

Conference Proceedings

C1. DUSK: Do not unlearn shared knowledge
Wonje Jeung*, Sangyeon Yoon*, Hyesoo Hong*, Soeun Kim, Seungju Han, Youngjae Yu, Albert No†
Findings of the Association for Computational Linguistics (**ACL Findings 2026)** [PDF]

C2. A2D: Any-Order, Any-Step Safety Alignment for Diffusion Language Models
Wonje Jeung*, Sangyeon Yoon*, Dongjae Jeon, Sangwoo Shin, Hyesoo Hong, Yoonjun Cho, Albert No†
International Conference on Learning Representations (**ICLR 2026**) [PDF]

C3. Rethinking Benign Relearning: Syntax as the Hidden Driver of Unlearning Failures
Sangyeon Yoon, Hyesoo Hong, **Wonje Jeung**, Albert No†
International Conference on Learning Representations (**ICLR 2026**)

C4. Rainbow Padding: Mitigating Early Termination in Instruction-Tuned Diffusion LLMs
BumJun Kim*, Dueun Kim*, Dongjae Jeon*, **Wonje Jeung**, Albert No†
International Conference on Learning Representations (**ICLR 2026**) [PDF]

C5. An Information Theoretic Evaluation Metric For Strong Unlearning
Dongjae Jeon*, **Wonje Jeung***, Taeheon Kim, Albert No†, Jonghyun Choi†
Association for the Advancement of Artificial Intelligence (**AAAI 2026**) [PDF]

C6. SAFEPATH: Preventing Harmfulness Reasoning in Chain-of-Thought via Early Alignment
Wonje Jeung, Sangyeon Yoon, Minsuk Kahng, Albert No†
Conference on Neural Information Processing Systems (**NeurIPS 2025**) [PDF]

C7. SEPS: A Separability Measure for Robust Unlearning in LLMs
Wonje Jeung*, Sangyeon Yoon*, Albert No†
Conference on Empirical Methods in Natural Language Processing (**EMNLP 2025**) [PDF]

C8. R-TOFU: Unlearning in Large Reasoning Models
Sangyeon Yoon, **Wonje Jeung**, Albert No†
Conference on Empirical Methods in Natural Language Processing (**EMNLP 2025**) [PDF]

C9. Large Language Models Still Exhibit Bias in Long Text
Wonje Jeung, Dongjae Jeon, Ashkan Yousefpour, Jonghyun Choi†
Findings of the Association for Computational Linguistics (**ACL Findings 2025**) [PDF]

C10. Representation Bending for Large Language Model Safety
Ashkan Yousefpour*, Taeheon Kim*, Ryan S. Kwon, Seungbeen Lee, **Wonje Jeung**, Seungju Han, Alvin Wan, Harrison Ngan, Youngjae Yu†, Jonghyun Choi†
Association for Computational Linguistics (**ACL 2025**) [PDF]

C11. ReALFRED: An Embodied Instruction Following Benchmark in Photo-Realistic Environments
Taewoong Kim*, Cheolhong Min*, Byeonghwi Kim, Jinyeon Kim, **Wonje Jeung**, Jonghyun Choi†
European Conference on Computer Vision (**ECCV 2024**) [PDF]

C12. Learning Equi-angular Representations for Online Continual Learning
Minhyuk Seo, Hyunseo Koh, **Wonje Jeung**, Minjae Lee, San Kim, Hankook Lee, Sungjun Cho, Sungik Choi, Hyunwoo Kim, Jonghyun Choi†
Conference on Computer Vision and Pattern Recognition (**CVPR 2024**) [PDF]

Workshops

W1. Adversarial Sample-Based Approach for Tighter Privacy Auditing in Final Model-Only Scenarios
Sangyeon Yoon*, **Wonje Jeung***, Albert No†
Workshop on Statistical Foundations of LLMs and Foundation Models (**NeurIPS SFLLM 2024**) [PDF]

W2. Multi-Level Knowledge Distillation and Dynamic Self-Supervised Learning for Continual Learning
SNUMPR TEAM
Workshop on CLVISION (**CVPR CLVISION 2024**), 2nd place [PDF]

Preprints

P1. Rewarding the Prompt, Not the Robot: Paraphrase-Induced Contradictions in Vision-Language Reward Models
Wonje Jeung, Sangyeon Yoon, Hyesoo Hong, Yoonjun Cho, Dongjae Jeon, Bumjun Kim, Youngjae Yu†, Jean Oh†, Albert No†
Under Review

P2. Be Helpful, Not Careful: Breaking Refusal Behavior with Minimal DPO Fine-Tuning
Sangyeon Yoon*, **Wonje Jeung***, Yoonjun Cho, Albert No†
Under Review

P3. BenchPreS: A Benchmark for Context-Aware Personalized Preference Selectivity of Persistent-Memory LLMs
Sangyeon Yoon, Sunkyoung Kim, Hyesoo Hong, **Wonje Jeung**, Yongil Kim, Wooseok Seo, Heuiyeen Yeen, Albert No†
Under Review [PDF]

Research Experience

AI-ISL (P.I. Albert No) — *Yonsei University*
Graduate Research Assistant — *Aug 2024 – Present*

- Showed that simply placing refusal responses in the DPO negative set bypasses OpenAI's moderation filter and breaks safety alignment in GPT-4o and GPT-4.1 (P2).
- Found that VLM-based evaluators for VLA models are surprisingly sensitive to prompt phrasing rather than actual task completion, and built a benchmark to capture this (P1).

- Developed safety alignment methods for LLMs (C10), LRMs (C6), and dLLMs (C2).
- Designed a privacy auditing algorithm with tight DP bounds for realistic black-box settings, using adversarially crafted samples (W1).
- Exposed limitations in existing unlearning evaluations and introduced an information-theoretic metric (C5), a separability measure (C7), and more realistic benchmarks (C8, C1).

Vnl Lab (P.I. Jonghyun Choi) — *Yonsei University*
Undergraduate Research Intern — *Mar 2023 – July 2024*

- Demonstrated that modern LLMs still exhibit systematic bias in long-form generation and proposed a mitigation strategy (C9).
- Ran large-scale annotation pipelines via Amazon MTurk to build an embodied instruction-following benchmark grounded in real-world environments (C11).
- Developed representation learning strategies for improved classification in online continual learning settings (W2, C12).

Professional Activities

Conference Reviewer

- ICLR, ICML, ICML Workshop (AIWILD) — *2026*
- T-IFS, NeurIPS Workshop (SoLaR) — *2025*

Invited Talks

- Gave a talk at the MLSYS group on deep safety alignment specialized for different language model architectures (LLMs, LRMs, dLLMs). — *Oct 2025*
- Presented at the MLSYS group on leveraging reasoning to achieve safety. — *June 2025*

Honors and Scholarship

James F. Lu World Fellows Program — *University of Michigan*
One of 4-6 recipients across all departments of the College of Engineering, with $5,000 merit stipend/year and $10,000 candidacy development stipend.

Y-Scholar Hub Research Excellence Selection — *Yonsei University*
Selected as one of seven for the Y-Scholar Hub research excellence profile.

Continual Learning Challenge — *CVPR 2024*
2nd Place Award.

Software Maestro Fellowship — IITP & ICT
Full Stipend ($6,154), elite software training program.

Academic Excellence Award — *Yonsei University*
1st semester 2020, 2nd semester 2020, 1st semester 2023, 1st semester 2024.

Academic Excellence Scholarship — *Yonsei University*
Covering full tuition fees, 2020.

Work Experience

SW Maestro — *Seoul, Korea*
Research Trainee, funded by Ministry of Science and ICT of Korean Government — *Apr 2022 – Dec 2022*

- Built a family app leveraging ML methods to enhance emotional connections, serving as team leader.
- Designed and implemented scalable AWS architecture and backend framework.

R2C Company — *Seoul, Korea*
Software Developer / Intern — *Oct 2021 – Apr 2022*

- Built a system to collect in-app user behavior data, enabling adaptive surveys based on user actions.
- Streamlined cross-platform development by unifying Android and iOS codebases.

Patents & Intellectual Property

Differential Privacy Auditing Framework *Jan 2025*
Computer program copyright (Science & Technology)

Skills

Research Pytorch, Huggingface, DeepSpeed, vLLM, Amazon MTurk Service
Development Spring, Node, Express, React, Flutter, Cloud Services (Lambda, S3, CloudFront, EC2)
Languages English: professional proficiency, Korean: native

Reference

Prof. Albert No
Associate Professor, Department of Artificial Intelligence, Yonsei University
Email: albertno@yonsei.ac.kr — Phone: +82 (2) 2123-7808
Prof. Minsuk Kahng
Assistant Professor, Department of Computer Science and Engineering, Yonsei University
Email: minsuk@yonsei.ac.kr
Prof. Jean Oh
Associate Professor, School of Computer Science (Robot Institute), Carnegie Mellon University
Email: hyaejino@andrew.cmu.edu