

Wonje Jeung

✉ Email 🎓 Scholar 🌐 Website in LinkedIn

Education

University of Michigan

PhD in Computer Science

Sept 2026- May 2030

(Expected)

- **James F. Lu World Fellows Program:** One of 4 recipients across all departments of the College of Engineering, with \$5,000 merit stipend/year and \$10,000 candidacy development stipend.
- Advisor: Lu Wang

Yonsei University

MS in Artificial Intelligence

Sept 2024 – Aug 2026

- GPA: 4.24/4.3
- **Coursework:** Information Theory, Statistical Pattern Recognition, Multimodal Deep Learning, Text Understanding, Large Scale Learning and Inference, Medical Imaging, Responsible AI
- Advisor: Albert No

BS in Computer Science

Mar 2020 – Aug 2024

- GPA: 4.09/4.3
- **Coursework:** Machine Learning, Computer Vision, Deep Learning Theory and Practice, Signal Processing, Probability and Statistics, Computer Systems, Computer Architecture, Operating Systems, Big Data
- Advisor: Jonghyun Choi

Publications

Conference Proceedings

- C1. Few-Shot Truly Benign DPO Attack for Jailbreaking LLMs
Sangyeon Yoon*, **Wonje Jeung***, Yoonjun Cho, Dongjae Jeon, Albert No[†]
Conference on Neural Information Processing Systems (NeurIPS 2026) [PDF] [🔗](#)
- C2. VLMS Trace Without Tracking: Diagnosing Failures in Visual Path Following
Hyesoo Hong, Minsoo Kim, **Wonje Jeung**, Sangyeon Yoon, Dongjae Jeon, Albert No[†]
Conference on Neural Information Processing Systems (NeurIPS 2026) [PDF] [🔗](#)
- C3. Systematic Hazard Sampling: Minimal-Variance Event Scheduling for CTMC/DTMC Discrete Diffusion and Flow Models
Seunghwan Jang, **Wonje Jeung**, Soojeon Han[†]
Conference on Neural Information Processing Systems (NeurIPS 2026)
- C4. Rewarding the Prompt, Not the Robot: Paraphrase-Induced Contradictions in Vision-Language Reward Models
Wonje Jeung, Sangyeon Yoon, Hyesoo Hong, Yoonjun Cho, Dongjae Jeon, Bumjun Kim, Youngjae Yu[†], Jean Oh[†], Albert No[†]
Conference on Empirical Methods in Natural Language Processing (EMNLP 2026)
- C5. BenchPreS: A Benchmark for Context-Aware Personalized Preference Selectivity of Persistent-Memory LLMs
Sangyeon Yoon, Sunkyoung Kim, Hyesoo Hong, **Wonje Jeung**, Yongil Kim, Wooseok Seo, Heuiyeon Yeen, Albert No[†]
Conference on Empirical Methods in Natural Language Processing (EMNLP 2026) [PDF] [🔗](#)
- C6. DUSK: Do not unlearn shared knowledge
Wonje Jeung*, Sangyeon Yoon*, Hyesoo Hong*, Soeun Kim, Seungju Han, Youngjae Yu, Albert No[†]
Findings of the Association for Computational Linguistics (ACL Findings 2026) [PDF] [🔗](#)
- C7. A2D: Any-Order, Any-Step Safety Alignment for Diffusion Language Models

Wonje Jeung*, Sangyeon Yoon*, Dongjae Jeon, Sangwoo Shin, Hyesoo Hong, Yoonjun Cho, Albert No[†]
International Conference on Learning Representations (ICLR 2026) [PDF] [↗](#)

- C8. Rethinking Benign Relearning: Syntax as the Hidden Driver of Unlearning Failures
Sangyeon Yoon, Hyesoo Hong, **Wonje Jeung**, Albert No[†]
International Conference on Learning Representations (ICLR 2026)
- C9. Rainbow Padding: Mitigating Early Termination in Instruction-Tuned Diffusion LLMs
BumJun Kim*, Dueun Kim*, Dongjae Jeon*, **Wonje Jeung**, Albert No[†]
International Conference on Learning Representations (ICLR 2026) [PDF] [↗](#)
- C10. An Information Theoretic Evaluation Metric For Strong Unlearning
Dongjae Jeon*, **Wonje Jeung***, Taeheon Kim, Albert No[†], Jonghyun Choi[†]
Association for the Advancement of Artificial Intelligence (AAAI 2026) [PDF] [↗](#)
- C11. SAFEPATH: Preventing Harmfulness Reasoning in Chain-of-Thought via Early Alignment
Wonje Jeung, Sangyeon Yoon, Minsuk Kahng, Albert No[†]
Conference on Neural Information Processing Systems (NeurIPS 2025) [PDF] [↗](#)
- C12. SEPS: A Separability Measure for Robust Unlearning in LLMs
Wonje Jeung*, Sangyeon Yoon*, Albert No[†]
Conference on Empirical Methods in Natural Language Processing (EMNLP 2025) [PDF] [↗](#)
- C13. R-TOFU: Unlearning in Large Reasoning Models
Sangyeon Yoon, **Wonje Jeung**, Albert No[†]
Conference on Empirical Methods in Natural Language Processing (EMNLP 2025) [PDF] [↗](#)
- C14. Large Language Models Still Exhibit Bias in Long Text
Wonje Jeung, Dongjae Jeon, Ashkan Yousefpour, Jonghyun Choi[†]
Findings of the Association for Computational Linguistics (ACL Findings 2025) [PDF] [↗](#)
- C15. Representation Bending for Large Language Model Safety
Ashkan Yousefpour*, Taeheon Kim*, Ryan S. Kwon, Seungbeen Lee, **Wonje Jeung**, Seungju Han, Alvin Wan, Harrison Ngan, Youngjae Yu[†], Jonghyun Choi[†]
Association for Computational Linguistics (ACL 2025) [PDF] [↗](#)
- C16. REALFRED: An Embodied Instruction Following Benchmark in Photo-Realistic Environments
Taewoong Kim*, Cheolhong Min*, Byeonghwi Kim, Jinyeon Kim, **Wonje Jeung**, Jonghyun Choi[†]
European Conference on Computer Vision (ECCV 2024) [PDF] [↗](#)
- C17. Learning Equi-angular Representations for Online Continual Learning
Minhyuk Seo, Hyunseo Koh, **Wonje Jeung**, Minjae Lee, San Kim, Hankook Lee, Sungjun Cho, Sungik Choi, Hyunwoo Kim, Jonghyun Choi[†]
Conference on Computer Vision and Pattern Recognition (CVPR 2024) [PDF] [↗](#)

Workshops

- W1. Adversarial Sample-Based Approach for Tighter Privacy Auditing in Final Model-Only Scenarios
Sangyeon Yoon*, **Wonje Jeung***, Albert No[†]
Workshop on Statistical Foundations of LLMs and Foundation Models (NeurIPS SFLLM 2024) [PDF] [↗](#)
- W2. Multi-Level Knowledge Distillation and Dynamic Self-Supervised Learning for Continual Learning
SNUMPR TEAM
Workshop on CLVISION (CVPR CLVISION 2024), 2nd place [PDF] [↗](#)

Research Experience

AI-ISL (P.I. Albert No)

Graduate Research Assistant

Yonsei University

Aug 2024 – Present

- Showed that simply placing refusal responses in the DPO negative set bypasses OpenAI’s moderation filter and breaks safety alignment in GPT-4o and GPT-4.1 (C1).
- Found that VLM-based evaluators for VLA models are surprisingly sensitive to prompt phrasing rather than actual task completion, and built a benchmark to capture this (??).
- Developed safety alignment methods for LLMs (C15), LRMs (C11), and dLLMs (C7).
- Designed a privacy auditing algorithm with tight DP bounds for realistic black-box settings, using adversarially crafted samples (W1).
- Exposed limitations in existing unlearning evaluations and introduced an information-theoretic metric (C10), a separability measure (C12), and more realistic benchmarks (C13, C6).

Vnl Lab (P.I. Jonghyun Choi)

Undergraduate Research Intern

Yonsei University

Mar 2023 – July 2024

- Demonstrated that modern LLMs still exhibit systematic bias in long-form generation and proposed a mitigation strategy (C14).
- Ran large-scale annotation pipelines via Amazon MTurk to build an embodied instruction-following benchmark grounded in real-world environments (C16).
- Developed representation learning strategies for improved classification in online continual learning settings (W2, C17).

Professional Activities

Conference Reviewer

- ICLR, ICLR Workshop (AIWILD), ICML, ICML Workshop (Mech Interp), 2026
NeurIPS
- T-IFS, NeurIPS Workshop (SoLaR) 2025

Invited Talks

- Gave a talk at the MLSYS group on deep safety alignment specialized for different language model architectures (LLMs, LRMs, dLLMs). Oct 2025
- Presented at the MLSYS group on leveraging reasoning to achieve safety. June 2025

Honors and Scholarship

James F. Lu World Fellows Program

One of 4-6 recipients across all departments of the College of Engineering, with \$5,000 merit stipend/year and \$10,000 candidacy development stipend.

University of Michigan

Y-Scholar Hub Research Excellence Selection

Selected as one of seven for the Y-Scholar Hub research excellence profile.

Yonsei University

Continual Learning Challenge

2nd Place Award.

CVPR 2024

Software Maestro Fellowship

Full Stipend (\$6,154), elite software training program.

IITP & ICT

Academic Excellence Award

1st semester 2020, 2nd semester 2020, 1st semester 2023, 1st semester 2024.

Yonsei University

Academic Excellence Scholarship

Covering full tuition fees, 2020.

Yonsei University

Work Experience

SW Maestro

Research Trainee, funded by Ministry of Science and ICT of Korean Government

Seoul, Korea

Apr 2022 – Dec 2022

- Built a family app leveraging ML methods to enhance emotional connections, serving as team leader.
- Designed and implemented scalable AWS architecture and backend framework.

R2C Company

Software Developer / Intern

Seoul, Korea

Oct 2021 – Apr 2022

- Built a system to collect in-app user behavior data, enabling adaptive surveys based on user actions.
- Streamlined cross-platform development by unifying Android and iOS codebases.

Patents & Intellectual Property

Differential Privacy Auditing Framework

Computer program copyright (Science & Technology)

Jan 2025

Skills

Research Pytorch, Huggingface, DeepSpeed, vLLM, Amazon MTurk Service

Development Spring, Node, Express, React, Flutter, Cloud Services (Lambda, S3, CloudFront, EC2)

Languages English: professional proficiency, Korean: native

Reference

Prof. Albert No

Associate Professor, Department of Artificial Intelligence, Yonsei University

Email: albertno@yonsei.ac.kr — Phone: +82 (2) 2123-7808

Prof. Minsuk Kahng

Assistant Professor, Department of Computer Science and Engineering, Yonsei University

Email: minsuk@yonsei.ac.kr

Prof. Jean Oh

Associate Professor, School of Computer Science (Robot Institute), Carnegie Mellon University

Email: hyaejino@andrew.cmu.edu